

ECE465/565 – Operating Systems Design: Project #2

Due date: October 16, 2023

Objectives

- To learn the concepts and methods related to process management such as process creation, process priorities, process scheduling and context switch.
- To understand Xinu's implementation of process management.
- To implement different scheduling algorithms in Xinu and verify their correct operation on representative test cases.

Reminders on Code Reuse, Testing and Use of ChatGPT.

Code reuse: This course has a **strict no-code reuse policy**: *consulting or reusing* (even partially) code from other individuals or online resources is not allowed (and is considered an academic integrity violation).

Testing: Your code will be graded using the course's VCL image. Please make sure you test your code on the VCL before submitting.

ChatGPT use: The use of ChatGPT (<https://chat.openai.com>) is allowed with the restrictions indicated in the syllabus (please consult the syllabus for guidelines). If you use ChatGPT, please list all the questions entered in it in the following form: <https://forms.gle/HgGTSVzF8CynFiCq5>.

Overview

This project focuses on process management. To help you proceed gradually, the project is divided into three parts.

In Part 1, your task is to learn Xinu's implementation of process management and perform some small coding tasks that will help you debug both Part 2 and Part 3. (**required for all students**)

In Part 2, you will implement [Lottery Scheduling \(LS\)](#) in Xinu. (**required for all students**)

In Part 3 (**required only for students taking the course at the graduate level – ECE565**), you will implement the [Multi-Level Feedback Queue \(MLFQ\)](#) scheduling policy in Xinu.

You should start from the unmodified Xinu's code. Before you proceed with your implementation, copy the whole `xinu` folder to a safe directory so to keep a clean version of Xinu. You should first implement Part 1. The code of Part 1 will be the baseline code used for Part 2 and Part 3. So, if you are taking the course at the graduate level, remember to save an extra copy of Part 1 before starting to code Part 2.

- If you are taking the course at the undergraduate level, you will submit a single code incorporating both Part 1 and Part 2.
- If you are taking the course at the graduate level, you will submit two versions of the code: (1) Part 1+2, and (2) Part 1+3.

Part 1: Understanding Xinu's process management and context switch.

The `xinu/system` folder contains the files related to process management and context switch. Study the files related to process creation (`create.c`), process scheduling (`resched.c`), context switch (`ctxsw.S`), process termination (`kill.c`), system initialization (`initialize.c`) and other related utilities (`ready.c`, `resume.c`, `suspend.c`, `chprio.c`, etc.).

Include in your report the answer to the following questions. Be clear and succinct.

- Q1. What is the `ready list`? List all the system calls that operate on it.
- Q2. What is the default process scheduling policy used in Xinu? Can this policy lead to process starvation? Explain your answer.
- Q3. When a context switch happens, how does the `resched` function decide if the currently executing process should be put back into the ready list?
- Q4. Analyze Xinu code and list all circumstances that can trigger a scheduling event.

Preliminary coding tasks:

CT1. Implement the function `syscall print_ready_list()` that prints the identifiers of the processes currently in the ready list. You don't need to document this function in your report.

CT2. Add to the PCB two fields:

- `uint32 runtime` – number of milliseconds the process has been running (i.e., in `PR_CURR` state).
- `uint32 turnaroundtime` – turnaround time in milliseconds.
- `uint32 num_ctxsw` – number of context switch operations *to* the process.

and update the code to keep these fields updated. Note that the `ctr1000` variable, which is meant to record the number of milliseconds elapsed since boot, is currently never updated by Xinu. Make sure you increment that variable in `clkhandler`.

CT3. In this project, we will group processes in two categories: *system processes* and *user processes*. In particular, we will call “system processes” the processes spawned by default by Xinu (e.g., `startup`, `main`, etc.), and “user processes” the ones used to model user applications. This second category includes the processes spawned by the `main` function. User processes can, in turn, spawn other user processes. For this project, you can ignore the `shell` process (as it won't be invoked in any of the test cases).

Write a `create_user_process` function to explicitly spawn user processes (this function can invoke the standard `create` function internally). The `create_user_process` function should have the following interface (note: it must not include a *priority* parameter).

```
pid32  create_user_process(  
    void      *funcaddr,          /* Address of the function */  
    uint32    ssize,              /* Stack size in bytes */  
    char      *name,              /* Name (for debugging) */  
    uint32    nargs,              /* Number of args that follow */  
    ...  
)
```

CT4. Instruments the `resched` function so that, when the `DEBUG_CTXSW` constant is defined, it will print all the context switch operations in the following format:

```
ctxsw::<old-process-id>-<new-process-id>
```

To this end, you can use the C preprocessor `#define` and `#ifdef/#endif` directives. Specifically, next to the `ctxsw` call you will have a block of code as below:

```
#ifdef DEBUG_CTXSW
    [statements to print the ctxsw information above]
#endif
```

In order to log the `ctxsw` calls, insert the `#define DEBUG_CTXSW` directive at the beginning of the `resched.c` file.

Note: You should log only real context switch operations (i.e., *the old and new process id must differ*).

For example, the following is a possible output for 3 user processes and the default Xinu scheduler with the `DEBUG_CTXSW` constant defined.

```
[...]
ctxsw::0-1
ctxsw::1-0
ctxsw::0-2
ctxsw::2-0
ctxsw::0-3
ctxsw::3-2
ctxsw::2-3
ctxsw::3-2
ctxsw::2-3
Obtained IP address 10.0.3.15 (0x0a00030f)
ctxsw::3-4
ctxsw::4-5
ctxsw::5-4
ctxsw::4-5
ctxsw::5-0
ctxsw::0-4
```

CT5. Write a function

```
void burst_execution(uint32 number_bursts, uint32 burst_duration,
                    uint32 sleep_duration);
```

which simulates the execution of applications that alternate execution phases requiring the CPU (CPU bursts), and execution phases not requiring it (CPU inactivity phases). Specifically:

- `number_bursts` = number of CPU bursts
- `burst_duration` = duration of each CPU burst in milliseconds (all CPU burst have the same duration)
- `sleep_duration` = duration of each CPU inactivity phase in milliseconds (all CPU inactivity phases have the same duration)

For consistency, include a CPU inactivity phase also at the end of the last CPU burst.

Note: burst and sleep duration model the phase behavior of the application. However, the OS scheduler may cause the application to wait for periods of time longer than the sleep duration.

Suggestion: have a look at the `timed_execution` function, which is provided in the test cases to simulate a CPU-intensive process with a predefined runtime.

Part 2: Lottery Scheduling

Your goal is to implement **Lottery Scheduling** in Xinu.

1. Implement the following function to initialize or modify the number of tickets (`tickets`) assigned to a given process (identified by `pid`).

```
void set_tickets(pid32 pid, uint32 tickets);
```

You can include the `set_ticket` function in the same file where you have implemented the `create_user_process` function, and invoke it from `main.c`. By default, a user process should be initialized with no tickets. *Note that this function is invoked from the test case files directly: you do not need to worry about ticket initialization and invoke it in your code.*

2. Utilize the `rand()` pseudo-random number generator from the `stdlib.h` library to generate the random numbers required by the scheduler. Use the default random seed (this is equivalent to invoking `srand(1)`).
3. Modify Xinu's scheduler so to allow lottery scheduling of the "user processes". Note that:
 - System processes must be scheduled with higher priority (using Xinu's default scheduling policy) and not follow the lottery scheduling policy. The *null* process should run only when there is nothing else that can be scheduled to run. You can test your scheduler by invoking user processes from the `main()` function. If you don't schedule the system processes with higher priority, you won't be able to test your lottery-based scheduler effectively.
 - For simplicity, **you can keep Xinu's scheduling events**, and not limit scheduling decisions to the end of each time slice.
 - The textbook (<http://pages.cs.wisc.edu/~remzi/OSTEP/cpu-sched-lottery.pdf>) provides a pseudocode for lottery scheduling. To make the implementation more efficient, sort the processes by descending number of tickets. If two processes have the same number of tickets, sort them by ascending process id (this is just a convention). In addition, your code should *run a lottery only when strictly required* (i.e., there are multiple user processes eligible to execute).
 - Recall that only processes in the ready list are eligible for execution. As a consequence, the tickets of processes that are not part of the ready list should not affect the scheduling decisions.
4. Analyze the fairness of your scheduler. To this end, write a test case file that spawns and runs two user processes with the same runtime, where the runtime is the execution time that the process would have if run from start to finish without preemption. Collect execution data with increasing runtime values. Plot the ratio between the turnaround times of the two processes when run together (i.e., $\text{turnaround-time}_{P1} / \text{turnaround-time}_{P2}$, where P1 and P2 are the two processes spawned) against the value of their runtime. You can select the number of data

points and the runtime values as you like, provided that your selection leads to a meaningful plot. Include the plot and a brief discussion of it (no more than a couple of sentences) in the report.

Include in the report:

- A *brief* description of your implementation approach, indicating the files involved in the implementation of lottery scheduling.
- The fairness analysis of (4) and the `main.c` file you have written in order to perform that analysis.

Be clear and succinct.

Part 3: Multi-Level Feedback Queue (MLFQ) scheduling policy – (only for ECE565 students).

Your goal is to implement the **Multi-Level Feedback Queue (MLFQ)** scheduling policy in Xinu. The MLFQ scheduling policy operates according to the following rules:

- *Rule 1:* if $\text{Priority}(A) > \text{Priority}(B)$, A runs
- *Rule 2:* if $\text{Priority}(A) = \text{Priority}(B)$, A&B run in RR fashion
- *Rule 3:* initially a job is placed at the highest priority level
- *Rule 4:* once a job uses up its time allotment, its priority is decreased
- *Rule 5:* after some time period S , move all jobs in the topmost queue

Implement the MLFQ scheduling policy in Xinu. Your implementation must follow the following directions:

- As in Part 2, system processes should be scheduled separately (using Xinu's default scheduling policy). In addition, you should keep Xinu's scheduling events, and not limit scheduling decisions to the end of each time slice. *In the absence of other scheduling events, you can let a process run until the end of a time slice even if it fully utilizes its time allotment within that time slice. Every time a process is moved to a different queue, its time allotment's utilization should be reset* (in other words, the scheduler does not need to keep track of situations where a process overuses its time allotment before its priority is reduced).
- The time allotment should double from one priority level to the next. For example, if the time allotment at priority p is T , the time allotment at priority $p-1$ should be $T*2$ (assume that p is higher than $p-1$).
- The number of priority levels (i.e., queues) for user processes (NUQ), the time allotment at the highest priority level (TA), and the priority boost period S should be configurable and defined in `include/resched.h` as follows:

```
#define UPRIORITY_QUEUES <NUQ>
#define TIME_ALLOTMENT <TA>
#define PRIORITY_BOOST_PERIOD <S>
```

Note: time allotment and priority boost period should be expressed in milliseconds.

- When performing priority boost, *first put the current process back to its queue*, and then enqueue the processes in order of priority. For example (in case of 3 priorities): if before priority boost $\text{HPQ} = \{1, 5\}$, $\text{MPQ} = \{2, 6\}$, $\text{LPQ} = \{4, 3\}$, after priority boost $\text{HPQ} = \{1, 5, 2, 6, 4, 3\}$ (HPQ, MPQ, LPQ being high, middle, and low priority queue, respectively). A new scheduling decision should be made every time there is a priority boost.
- Add to the PCB fields `upgrades` and `downgrades` keeping track of the number of priority upgrades and downgrades undergone by each process.

- Add a `void reset_timing()` function that can be called from the test case file to reset the priority upgrade period. For example, if you use `ctr1000` and another counter to determine when priority upgrades need to be triggered, the function will look something like this:

```
void reset_timing(){
    ctr1000 = 0;
    /* reset counter used to trigger priority upgrades */
}
```

Include in the report: a succinct description of your implementation approach, indicating the files involved in the implementation of the MLFQ scheduling policy.

Test case files

You can use the test case files provided to debug your code and verify its correct operation. The test case files are organized in two folders: *lottery* (test cases for lottery scheduling) and *mlfq* (test cases for MLFQ). Each folder contains a README file, *including important information on the settings used in the test cases and their structure*. Please make sure you read the README files and use the settings indicated. Note that the outcome of the test cases might change slightly from run to run, and small deviations from the reference outputs are expected. Being able to identify acceptable outputs based on the operation of lottery scheduling and of MLFQ is part of this exercise. The README files provide some guidance.

Submission instructions

1. **Important:** We will test your implementations using different test cases (i.e., different `main.c` files). Therefore, do not implement any essential functionality (other than your test cases) in the `main.c` file. Also, turn off debugging output before submitting your code.
2. **Testing:** Remember to test your code using the course's VCL image (*Xinu (ECE565)*).
3. Go to the `xinu/compile` directory and invoke `make clean`.
4. As for the previous project, create a `xinu/tmp` folder and **copy** all the files you have modified/created into it (the `tmp` folder should have the same directory structure as the `xinu` folder).
5. Go to the parent folder of the `xinu` folder. You will submit Part 2 and Part 3 separately (the code of Part 1 should be incorporated in both Part 2 and Part 3). For each part (2 & 3), compress the whole `xinu` directory into a `tgz` file.

```
tar czf xinu_project2_part#.tgz xinu
```

6. Submit your assignment – including `tgz` files and report – through Moodle. Please upload only one `tgz` file **for each Part**. There is no need to print the report and bring it to class.
7. Please save the report in pdf format, and submit it separately (i.e., please do not include the report in the code archives).

Grading rubric

		ECE465 students	ECE565 students
<i>Part 1</i>	<i>Q1-Q4</i>	8	5
	<i>CT1</i>	5	3
	<i>CT2</i>	3	2
	<i>CT3</i>	1	1
	<i>CT4</i>	1	1
	<i>CT5</i>	5	3
<i>Part 2</i>	<i>Implementation of lottery scheduling</i>	55 (20 for good effort)	35 (10 for good effort)
	<i>Description of the implementation</i>	10	5
	<i>Analysis of fairness</i>	12	5
<i>Part 3</i>	<i>Implementation of MLFQ</i>		35 (10 for good effort)
	<i>Description of implementation</i>		5